Does the Quantum Layer Matter? A Controlled Study of Quantum-Assisted Physics-Informed Neural Networks on Burgers' Equation

Hemanth Bharatula
Independent researcher
EMAIL-TO-CONFIRM

October 2026

Abstract

Quantum-assisted physics-informed neural networks (QAPINNs) replace the first hidden layer of a PINN with a variational quantum circuit (VQC) and are reported to match classical accuracy with fewer trainable parameters. We ask whether the quantum layer itself is responsible for any change in learning, using a verified, parameter-matched bench on 1-D viscous Burgers' equation with an exact Cole–Hopf reference, converged training (Adam followed by L-BFGS), up to six seeds per configuration (ten for classical models), and four kinds of classical control: parameter-matched, structural (the same post-network fed by a trainable classical front end), depth-matched, and random Fourier features with the VQC's frequency range. At $\sim$121 parameters the classical PINN has the lowest error (0.270 ± 0.041) and is significantly better than both depth-3 QAPINN variants (Welch $p \leq 0.001$). At $\sim$1273 parameters the best QAPINN (0.094 ± 0.082) beats shallow and balanced classical networks ($p \leq 0.007$) but is statistically indistinguishable from a depth-matched classical network (0.104 ± 0.077, $p = 0.82$) and from the structural control (0.138 ± 0.084, $p = 0.34$). Entanglement has no detectable effect at either scale ($p = 0.42$, $p = 0.44$). The quantum layer is not a fixed feature map: four *fixed* random Fourier features fail badly (0.73–0.85), but a trainable classical front end of the same width matches the VQC, and sixteen fixed features tie it. A spectral analysis of 111 trained models shows that, among models that fit the smooth part of the solution, error tracks the realised output bandwidth (Spearman $\rho = -0.91$ at small scale): small QAPINNs stay near the low band of their features, and a larger classical post-network extends any front end's band to the reference. We find no quantum-specific advantage on this problem, and we distil the checks that removed three of our own initial positive findings into a four-step protocol for evaluating such claims.

1 Introduction

Physics-informed neural networks (PINNs) [1] solve partial differential equations by training a network $u_\theta(x, t)$ to minimise the PDE residual together with initial and boundary conditions. A quantum-assisted PINN (QAPINN) [2] replaces the first hidden layer with a VQC. On 1-D Burgers' equation this has been reported to keep accuracy while using about 20% fewer trainable parameters [2], and related hybrid designs report similar gains [3]. Such comparisons are usually single-seed, use one classical baseline of a different shape, and are not compute-matched, so they cannot separate the effect of the quantum layer from that of architecture, optimisation or seed.

Recent work on benchmarking quantum machine learning [18] and on dequantising variational models with classical Fourier features [15, 16, 17] suggests that such effects should be tested against strong classical baselines; to our knowledge this has not been done for physics-informed training, where the model is fitted to a PDE residual rather than to labelled data.

The standard theoretical lens for VQCs is that a circuit with data re-uploading is a truncated Fourier series in its inputs [4, 5]. We ask whether this lens, together with properly controlled experiments, explains what the quantum layer does inside a PINN: *when* it changes learning, *why*, and whether any change is specific to the quantum component.

Contributions.

1. A verified, parameter-matched bench for Burgers' and heat equations with exact references, an explainability suite (spectra, neurons, VQC internals, loss landscape, generalisation), and released code that regenerates every number in this paper from the run logs.
2. A converged, multi-seed comparison at two model scales against parameter-matched, structural, depth-matched and random-Fourier-feature classical controls, which finds no statistically significant quantum-specific advantage and no detectable entanglement effect.
3. A mechanistic account: the VQC acts as a *trainable* low-frequency feature map that a trainable classical front end of the same width reproduces; across 111 trained models, error tracks the realised output bandwidth, which the classical post-network, not the quantum layer, determines.
4. A four-step protocol (converge, use fair controls, require multi-seed significance, separate features from model) that removed three of our own initial positive findings.

2 Background and related work

VQCs as Fourier series. A VQC whose inputs enter through Pauli-rotation encodings computes a partial Fourier series $f(x) = \sum_{\omega \in \Omega} c_\omega e^{i\omega x}$ [4]. The frequency set Ω is fixed by the encoding; the trainable gates set the coefficients. With L re-uploads of a single-qubit R_Y encoding, $\Omega = \{-L, \ldots, L\}$, so the bandwidth grows with depth [4, 5]. Mhiri et al. [6] show that a frequency being available does not mean its coefficient is reachable: coefficient variance depends on frequency redundancy and can vanish with qubit count. Trainability of VQCs is further limited by barren plateaus [7].

Spectral bias in classical networks. Gradient-trained networks fit low frequencies first (the frequency principle or spectral bias) [8, 9], which is one reason PINNs struggle with sharp features [10, 11]. Classical Fourier-feature front ends were introduced to counter this bias [12]; a VQC front end is structurally similar, which motivates our classical front-end control.

Dequantisation and benchmarking. Because VQC outputs are Fourier series, a classical model built on random Fourier features drawn from the circuit's frequency spectrum can approximate many variational models [15]; the conditions under which this works, and where it fails, have since been analysed for regression and classification [16, 17]. Bowles et al. [18] show in a large benchmark that out-of-the-box classical models usually match or beat quantum classifiers, and that results are sensitive to hyperparameters and baselines. These studies concern supervised learning; we test the same questions for a VQC embedded in a PINN and trained only through a PDE residual.

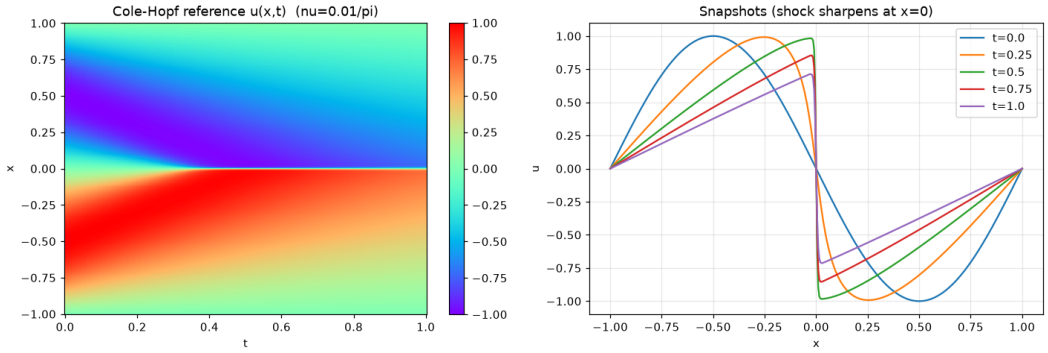


Figure 1: Cole–Hopf reference for viscous Burgers. Left: $u(x,t)$; right: snapshots showing the front forming at $x = 0$.

Quantum PDE solvers. Kyriienko et al. [13] solve nonlinear differential equations with differentiable quantum circuits. Hybrid quantum-classical PINNs [2, 3] report parameter savings at similar accuracy, with PDE-dependent results (quantum variants sometimes lose, e.g. on wave and cavity problems [3]), and a hybrid quantum PINN for high-speed flows was found to be efficient on harmonic problems but to struggle on non-harmonic ones [19]. Quantum PINNs can also collapse to trivial solutions [14]. Mechanistic explanations in this literature usually appeal to entanglement-driven expressivity without a controlled test, which is the gap we target.

3 Setup

3.1 Problems

Both PDEs use $x \in [-1, 1]$, $t \in [0, 1]$, $u(x, 0) = -\sin(\pi x)$ and homogeneous Dirichlet boundaries, so only the operator differs:

$$\text{Burgers:} \quad u_t + u\,u_x = \nu\,u_{xx}, \quad \nu = 0.01/\pi, \tag{1}$$

$$\text{Heat:} \quad u_t = \alpha\,u_{xx}, \quad \alpha = 1/\pi^2, \quad u = -\sin(\pi x)e^{-t}. \tag{2}$$

Burgers develops a steep front at $x = 0$ (Fig. 1); its reference is the Cole–Hopf solution evaluated by Gauss–Hermite quadrature. Heat has a single spatial mode. All results in the main text are on Burgers; heat serves as a sanity check. Error is the relative L_2 norm on a 256×100 grid.

3.2 Models

QAPINN: $(x, t) \to$ VQC $\to$ classical post-network $\to u$. A fixed map sends coordinate $i \bmod 2$ to qubit i (no trainable pre-network, so the Fourier band lives directly in (x, t)). The VQC has 4 qubits, L layers of R_Y angle encoding with re-uploading, trainable rotations, a CNOT pattern (*linear* chain or *none*), and $\langle Z_i \rangle$ readout. The small post-network is $[4, 16, 1]$ ($L = 3$: 24 quantum + 97 classical = 121 parameters); the larger one is $[4, 32, 32, 1]$ (1273 parameters). All networks use tanh.

Controls (Table 1): a *parameter-matched* classical MLP; a *structural* control (CFE) that keeps the QAPINN's post-network but replaces the VQC with a classical Linear$(2 \to 4)$ + tanh front end, so the only difference is a quantum versus classical front end; and, at the larger scale, a *depth-matched* classical MLP with the same number of hidden layers as the QAPINN counting the VQC

Scale	Model	Architecture	Params
Small	QAPINN, depth 3 (linear / none)	VQC(4q, L=3) $\to [4, 16, 1]$	121
	QAPINN, depth 5 (linear)	VQC(4q, L=5) $\to [4, 16, 1]$	137
	Classical	$[2, 7, 11, 1]$	121
	CFE (structural)	Linear(2$\to$4)+tanh $\to [4, 16, 1]$	109
	RFF, width 4 / 16	fixed cos features $\to [\cdot, 16, 1]$	97 / 289
Larger	QAPINN, depth 3 (linear / none)	VQC(4q, L=3) $\to [4, 32, 32, 1]$	1273
	CFE (structural)	Linear(2$\to$4)+tanh $\to [4, 32, 32, 1]$	1261
	Classical, depth-matched	$[2, 22, 22, 22, 1]$	1101
	Classical, balanced	$[2, 32, 32, 1]$	1185
	Classical, shallow	$[2, 16, 68, 1]$	1273
	RFF, width 4 / 16	fixed cos features $\to [\cdot, 32, 32, 1]$	1249 / 1633

Table 1: Models in the converged study. Small: Adam 2500 + L-BFGS 400 iterations, 4000 interior points. Larger: Adam 3000 + L-BFGS 600, 6000 interior points. Both use learning rate 10^{-2} and 150 boundary/initial points, identical for every model.

as one, plus a *balanced* two-layer MLP. Finally, a *random-Fourier-feature* control (RFF) replaces the VQC with fixed, untrained features $\cos(w_k \cdot (x, t) + b_k)$ whose integer frequencies $w_k \in \{-6, \dots, 6\}^2$ cover the VQC's measured range (Section 4.5) and whose phases are uniform, drawn afresh for each seed, feeding the identical post-network. We use width 4 (same post-network input as the QAPINN, fewer parameters) and width 16 (more features, more parameters).

3.3 Training, statistics and verification

Every model at a given scale is trained under the identical protocol of Table 1. Seeds $0, \dots, n-1$ are used: $n = 10$ for the classical baselines and the trainable-front-end control, $n = 6$ for QAPINNs and the RFF controls, and $n = 3$ for the depth-5 QAPINN. We report mean $\pm$ sample standard deviation and two-sided Welch t-test p-values (denoted p), computed from the per-run logs by a script in the repository; Student and Mann–Whitney values are given there for comparison. The VQC is simulated with PennyLane [20] (`default.qubit`, exact statevector, backpropagation for the second derivatives the residual needs).

Before any comparison we checked every load-bearing computation against an independent reference (Appendix A): autograd residuals against analytic ones, derivatives against finite differences, the error metric on the exact solution, parameter counts, and same-seed reproducibility. A one-qubit sweep confirms that the VQC output contains no harmonic above the re-upload depth for $L = 1, \dots, 6$.

4 Results

4.1 Small scale: the classical PINN is best

Table 2 and Fig. 2 (left) give the converged results at $\sim$121 parameters. The classical MLP has the lowest mean error and is significantly better than both depth-3 QAPINNs ($p = 0.001$ against the linear-entanglement reference design, $p = 0.0001$ against the unentangled variant). Increasing VQC depth to 5 lowers the QAPINN's error to the level of the structural control ($p = 0.91$); the gap to the classical MLP is then not significant ($p = 0.49$, $n = 3$).

Model	Params	relative L_2 (mean $\pm$ sd)	n
Classical [2, 7, 11, 1]	121	$\mathbf{0.270 \pm 0.041}$	10
QAPINN depth 5, linear	137	0.314 ± 0.090	3
CFE (structural)	109	0.321 ± 0.054	10
QAPINN depth 3, none	121	0.379 ± 0.035	6
QAPINN depth 3, linear	121	0.402 ± 0.056	6
RFF, width 16	289	0.416 ± 0.016	6
RFF, width 4	97	0.847 ± 0.086	6

Table 2: Converged results at small scale on Burgers.

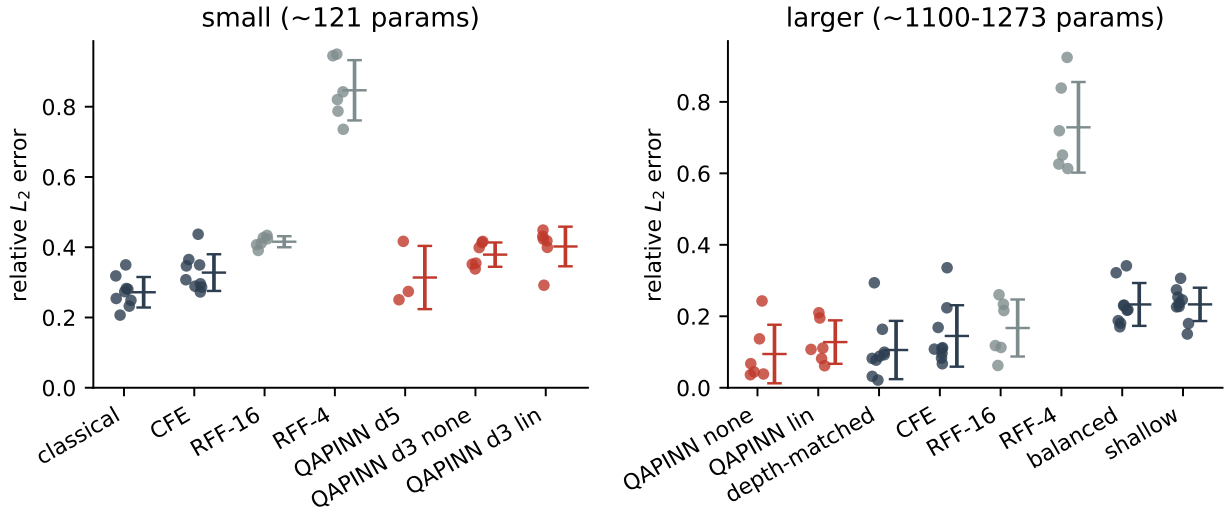


Figure 2: Per-seed relative L_2 error (dots) with mean $\pm$ sd (bars). Red: QAPINN; dark: classical models and controls; grey: random-Fourier-feature controls. Left: small scale. Right: larger scale; the best QAPINN is separated from the shallow and balanced classical networks but not from the depth-matched or structural controls, and seed-to-seed spread is large.

4.2 Larger scale: the apparent advantage is architectural

At the small scale every model still leaves the front partly unresolved (relative L_2 0.25–0.40). We therefore repeated the study at $\sim$1273 parameters, where the best runs reach relative $L_2 \approx 0.03$–0.04 (Table 3, Fig. 2, right).

Against the shallow and balanced classical networks the QAPINN appears clearly better ($p = 0.005$ and $p = 0.007$, both surviving a Holm correction over the seven comparisons). Against the fair controls it is not: a depth-matched classical MLP with *fewer* parameters has almost the same mean error, and the structural control is also statistically indistinguishable from it. The advantage over the shallow network is therefore attributable to depth and post-network shape, which a classical model reproduces, rather than to the quantum layer.

This is a failure to reject, not a demonstration of equality. Seed-to-seed spread (sd ≈ 0.08) exceeds the mean differences (≈ 0.01–0.04), so the comparison is under-powered. One QAPINN-none seed (0.243) is an outlier; removing it lowers the mean but the comparison with the depth-matched control remains non-significant ($p = 0.22$).

Model	Params	n	relative L_2 (mean $\pm$ sd)	vs QAPINN-none
QAPINN depth 3, none	1273	6	$\mathbf{0.094 \pm 0.082}$	—
Classical, depth-matched	1101	10	0.104 ± 0.077	$p = 0.82$
QAPINN depth 3, linear	1273	6	0.128 ± 0.061	$p = 0.44$
CFE (structural)	1261	10	0.138 ± 0.084	$p = 0.34$
RFF, width 16	1633	6	0.167 ± 0.080	$p = 0.15$
Classical, balanced	1185	10	0.231 ± 0.057	$p = 0.007$
Classical, shallow	1273	10	0.242 ± 0.052	$p = 0.005$
RFF, width 4	1249	6	0.729 ± 0.127	$p < 10^{-4}$

Table 3: Converged results at larger scale, ordered by mean error. p: two-sided Welch t-test against QAPINN-none.

4.3 Entanglement has no detectable effect

With three seeds at small scale, the unentangled QAPINN looked better than the linear one (0.369 vs 0.424; Student $p = 0.024$, Welch $p = 0.060$). Extending to six seeds removed the effect (0.379 vs 0.402, $p = 0.42$); one of the new linear seeds (0.292) is the best depth-3 run. At the larger scale the difference is also not significant (0.094 vs 0.128, $p = 0.44$). Earlier single-seed ablations with four entanglement patterns (Appendix B) likewise showed entanglement entropy varying widely with no corresponding change in error.

4.4 The quantum layer is a trainable, not a fixed, feature map

Because VQC outputs are Fourier series, the natural classical surrogate is a set of random Fourier features from the same frequency range [15, 16]. With four fixed features (the QAPINN's post-network input width), this surrogate fails at both scales (0.847 and 0.729; $p < 10^{-4}$ against the QAPINN): four randomly chosen cosines do not span even the smooth part of the solution, and nothing in front of the post-network can adjust them. The VQC's 24 trainable rotations do adjust its four features, and that matters. But a *trainable* classical front end of the same width (the structural control, 12 front-end parameters) matches the VQC: it has lower mean error at small scale and is statistically indistinguishable at larger scale. And sixteen fixed features tie the QAPINN without any training of the front end (small scale 0.416 vs 0.402, $p = 0.59$; larger scale 0.167 vs 0.094, $p = 0.15$, and vs 0.128 for the entangled QAPINN, $p = 0.36$). The quantum layer is therefore reproducible classically, either by training a narrow front end or by widening a fixed one.

4.5 Error tracks output bandwidth

For every trained model we computed the spatial spectrum of u_θ on the evaluation grid (magnitude of the FFT along x, averaged over t) and its 99%-energy bandwidth K_{99}: the smallest K such that harmonics $1, \ldots, K$ hold 99% of the non-constant spectral energy. The reference has $K_{99} = 13$.

Figure 3 shows one trend across model families. Excluding the width-4 RFF models, which fail to fit even the smooth part of the solution, error falls with bandwidth (Spearman $\rho = -0.81$, $n = 99$). The relation is strongest at small scale ($\rho = -0.91$, $n = 41$; $\rho = -0.94$ within the 15 small QAPINNs): depth-3 QAPINNs mostly realise $K_{99} = 2$–6 and the 16-feature RFF control 2–4, both with error near 0.4, while the small classical MLPs reach 6–10. The one depth-3 seed that reaches $K_{99} = 7$ is also the most accurate. At the larger scale every QAPINN, structural control and depth-matched network reaches $K_{99} = 11$–14 (some shallow, balanced and RFF runs stay at 6–8). Bandwidth then explains much less of the error ($\rho = -0.35$, $n = 58$): models at the reference

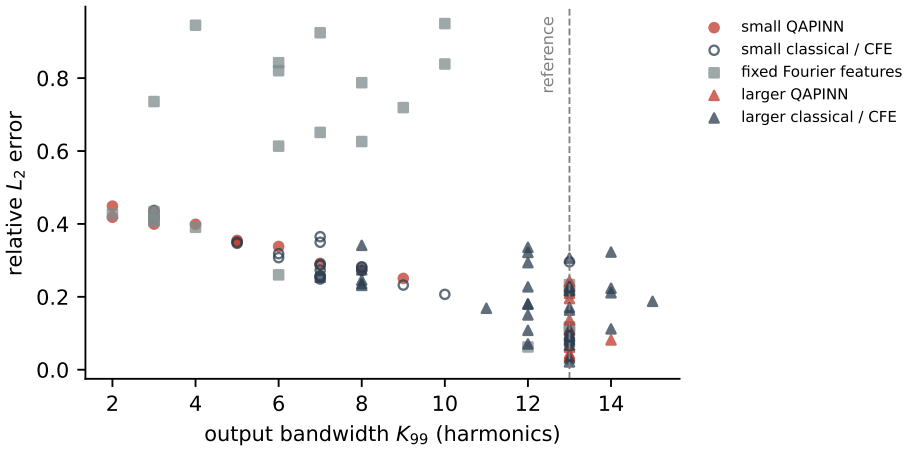


Figure 3: Relative L_2 error against the realised output bandwidth K_{99} for all 111 converged models (circles: small scale; triangles: larger scale; red: QAPINN; grey squares: random-Fourier-feature controls). The dashed line is the reference bandwidth. The width-4 RFF models (top) fail for a different reason (Section 4.4).

bandwidth still range from 0.03 to 0.34, a spread set by how well each seed's optimisation resolves the front rather than by which frequencies are present. Bandwidth is thus necessary for low error but not sufficient.

This corrects a claim in our own earlier analysis. The VQC's *features* are band-limited by construction. With $R_Y(x)$ encoding, $L = 3$ and two of the four qubits encoding x, each readout of a trained unentangled QAPINN contains angular frequencies in x up to 3; with the linear CNOT chain, which mixes the two x-qubits, the readouts reach 4–6 (measured exactly, since the features are 2π-periodic in x). Either way this is at most about two periods across $[-1, 1]$, far below the 13 harmonics the reference needs. The *model* is not: the tanh post-network forms products of features and so can synthesise higher harmonics. Whether it does depends on its capacity. The small post-network mostly fails to; the larger one always succeeds. The band set by VQC depth is therefore a soft constraint on small models, not a property of the QAPINN as such.

4.6 Encoding scale

In a single-seed preliminary sweep (Appendix B) the input scale of the encoding had an optimum near 1. At scale 2π the model collapsed to the trivial solution $u \equiv 0$ (relative $L_2 = 1.000$), which satisfies the boundary conditions but not the initial condition. With scale 2π each re-upload wraps the domain twice, so nearby collocation points map to distant circuit states; we suggest this as the cause but did not test it further.

5 Interpretation

We mark each explanation as *measured* (directly supported by our data) or *reasoned* (consistent with theory and our controls but not isolated).

The quantum layer is a trainable low-frequency feature map (measured). With raw coordinates as input and $L = 3$, the VQC supplies the post-network with trigonometric features

whose frequencies are fixed by the encoding and whose coefficients are trained. The training matters: four fixed features of the same frequency range fail (Section 4.4). Error then tracks how far the post-network can extend the features' band (Section 4.5). A trainable classical front end of the same width gives the post-network features that are as useful, and sixteen fixed Fourier features are enough to match it.

Why there is no quantum-specific advantage (reasoned). The representational work is done by the classical post-network, which sits downstream of the quantum layer. A front end only needs to supply features from which the post-network can build the solution. Burgers' solution is low-frequency away from the front and needs many harmonics at the front; the VQC's band offers no shortcut to the latter, so any front end of comparable richness performs similarly. This mirrors the dequantisation results for supervised learning [15, 16]: a model whose quantum part is a trainable Fourier series can be matched by a classical model with access to similar features.

Why entanglement is neutral (partly measured). Entanglement does change the features: the linear CNOT chain doubles the highest x-frequency of the readouts from 3 to 6 (Section 4.5) and redistributes which coefficients are reachable [6]. Yet it has no detectable effect on error. Our interpretation is that both changes are redundant once a post-network follows: products of the unentangled per-qubit features already contain these sum frequencies, and the post-network forms such products itself. We tried to measure the mechanism directly (per-qubit feature correlation and effective rank, with and without entanglement) and the two measures disagreed, so we report this as an interpretation consistent with the null result.

Why the front defeats small models (partly measured). The front's Fourier content decays slowly with k, and gradient-trained networks fit high frequencies last [8, 9]. For the larger QAPINNs where training-time spectra were logged, the order in which harmonics reach half their initial error correlates positively with k in 6 of 8 runs (Spearman ρ from -0.16 to 0.84), weak support for this classical explanation. The quantum layer neither causes nor removes it.

6 A protocol for evaluating QAPINN claims

Three positive findings appeared during this project and were removed by later checks: (i) a larger diversity of post-network neurons with a quantum front end, which disappeared against the structural control (Appendix B); (ii) an entanglement effect at $n = 3$, which disappeared at $n = 6$ (Section 4.3); and (iii) an advantage at larger scale, which disappeared against a depth-matched classical network (Section 4.2). A fourth reading, that VQC depth caps the model at a fixed low band, holds for the features and loosely for small models, but not for models with a larger post-network (Section 4.5). Each of the first three would have been reported as positive without the corresponding check:

1. **Converge.** Train with a second-order optimiser (here L-BFGS) after Adam. Without it, every Burgers model in our preliminary runs stalled at relative L_2 0.42–0.59 and the differences between them were within seed noise.
2. **Use fair controls.** Compare against a parameter-matched, a structurally matched (trainable classical front end) and a depth-matched classical network, plus a Fourier-feature model with the circuit's frequency range, all trained under the same protocol.
3. **Require multi-seed significance.** Use at least six seeds (more for cheap classical models), report the test used, and treat $n = 3$ effects as provisional.

4. **Separate features from model.** Test Fourier-band claims on the VQC output and on the trained model separately.

An advantage that passes all four checks would justify a circuit-design methodology. None of ours did.

7 Limitations

1. **Statistical power.** With $n = 6$ for the QAPINNs and seed spread of $\sim$0.08, the decisive comparisons are failures to reject. Same-seed reruns that include L-BFGS are not bit-identical across machines, which adds variance. More QAPINN seeds would tighten the null results; each costs hours of simulation.
2. **Scope.** One nonlinear PDE in one spatial dimension (plus heat as a sanity check), two model scales, 4-qubit circuits, and R_Y encoding of raw coordinates. All models share one learning rate and schedule; per-model tuning could change individual results, though it would not obviously favour the quantum model. A lid-driven cavity case is implemented in the code but was not run. Other encodings, larger circuits or other PDEs could behave differently.
3. **Simulation only.** Exact statevector simulation without shot noise or hardware noise, which would only add difficulties for the quantum model. Fewer parameters do not mean lower cost: QAPINN training took roughly 130–500 times longer in wall-clock time than the classical models (CPU laptops).
4. **Untested mechanisms.** The explanations marked *reasoned* in Section 5 were not isolated experimentally, and the encoding-scale collapse is a single-seed observation.

8 Conclusion

On 1-D Burgers' equation, with converged training, multi-seed statistics and fair classical controls, we find no quantum-specific advantage for QAPINNs: at small scale the classical PINN is best, and at larger scale the QAPINN ties both a depth-matched classical network and a structural control. Entanglement has no detectable effect. The VQC acts as a trainable low-frequency feature map: a trainable classical front end of the same width, or a wider set of fixed Fourier features, matches it. Across 111 trained models, error tracks the bandwidth the model realises, which the classical post-network, not the quantum layer, determines. Whether an advantage appears on other PDEs, encodings or scales remains open; the protocol in Section 6 is a way to test such a claim.

Code and data availability

Code, configurations, per-run logs and metrics, and the trained weights of the converged models are available at
`https://github.com/hementhu/qapinn-bqp`. The scripts `compute_stats.py`, `verify_claims.py` and `verify_pipeline.py` regenerate every statistic, spectral measurement and verification check reported here.

Acknowledgements

This work began as a submission to the BQP industry challenge of the WISER Global Quantum+AI Program 2026. The author used Claude Code (Anthropic) to help implement the code, run and

aggregate experiments, search the literature, audit results and draft text. The author directed the research questions and experiments, checked the results, and takes responsibility for the content.

References

[1] M. Raissi, P. Perdikaris, G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *J. Comput. Phys.* **378**, 686–707 (2019).

[2] J. Shah, R. Lineswala, A. Chopra. Benchmarking quantum-assisted physics-informed neural networks (QA-PINN) for computational fluid dynamics. *IEEE International Conference on Quantum Computing and Engineering (QCE)* (2024). [VENUE/DOI TO CONFIRM]

[3] A. Farea, S. Khan, M. S. Celebi. QCPINN: Quantum-classical physics-informed neural networks for solving PDEs. arXiv:2503.16678 (2025).

[4] M. Schuld, R. Sweke, J. J. Meyer. Effect of data encoding on the expressive power of variational quantum-machine-learning models. *Phys. Rev. A* **103**, 032430 (2021).

[5] A. Pérez-Salinas, A. Cervera-Lierta, E. Gil-Fuster, J. I. Latorre. Data re-uploading for a universal quantum classifier. *Quantum* **4**, 226 (2020).

[6] H. Mhiri, L. Monbroussou, M. Herrero-González, S. Thabet, E. Kashefi, J. Landman. Constrained and vanishing expressivity of quantum Fourier models. *Quantum* **9**, 1847 (2025).

[7] J. R. McClean, S. Boixo, V. N. Smelyanskiy, R. Babbush, H. Neven. Barren plateaus in quantum neural network training landscapes. *Nat. Commun.* **9**, 4812 (2018).

[8] Z.-Q. J. Xu, Y. Zhang, T. Luo, Y. Xiao, Z. Ma. Frequency principle: Fourier analysis sheds light on deep neural networks. *Commun. Comput. Phys.* **28**, 1746–1767 (2020).

[9] N. Rahaman, A. Baratin, D. Arpit, F. Draxler, M. Lin, F. Hamprecht, Y. Bengio, A. Courville. On the spectral bias of neural networks. *ICML* (2019).

[10] S. Wang, Y. Teng, P. Perdikaris. Understanding and mitigating gradient flow pathologies in physics-informed neural networks. *SIAM J. Sci. Comput.* **43**, A3055–A3081 (2021).

[11] A. S. Krishnapriyan, A. Gholami, S. Zhe, R. M. Kirby, M. W. Mahoney. Characterizing possible failure modes in physics-informed neural networks. *NeurIPS* (2021).

[12] M. Tancik, P. P. Srinivasan, B. Mildenhall, S. Fridovich-Keil, N. Raghavan, U. Singhal, R. Ramamoorthi, J. T. Barron, R. Ng. Fourier features let networks learn high frequency functions in low dimensional domains. *NeurIPS* (2020).

[13] O. Kyriienko, A. E. Paine, V. E. Elfving. Solving nonlinear differential equations with differentiable quantum circuits. *Phys. Rev. A* **103**, 052416 (2021).

[14] Chen, Shaviner, Chandravamsi, Pisnoy, Frankel, Pereg. Quantum physics-informed neural networks for Maxwell's equations. arXiv:2506.23246 (2025).

[15] J. Landman, S. Thabet, C. Dalyac, H. Mhiri, E. Kashefi. Classically approximating variational quantum machine learning with random Fourier features. arXiv:2210.13200 (2022).

[16] R. Sweke, E. Recio-Armengol, S. Jerbi, E. Gil-Fuster, B. Fuller, J. Eisert, J. J. Meyer. Potential and limitations of random Fourier features for dequantizing quantum machine learning. *Quantum* **9**, 1640 (2025).

[17] M. Sahebi, A. Barthe, Y. Suzuki, Z. Holmes, M. Grossi. On dequantization of supervised quantum machine learning via random Fourier features. arXiv:2505.15902 (2025).

[18] J. Bowles, S. Ahmed, M. Schuld. Better than classical? The subtle art of benchmarking quantum machine learning models. arXiv:2403.07059 (2024).

[19] F. Y. Leong, W.-B. Ewe, T. S. B. Quang, Z. Zhang, J. Y. Khoo. Hybrid quantum physics-informed neural network: Towards efficient learning of high-speed flows. arXiv:2503.02202 (2025).

[20] V. Bergholm et al. PennyLane: Automatic differentiation of hybrid quantum-classical computations. arXiv:1811.04968 (2018).

A Pipeline verification

Check	Method	Result
Burgers residual	autograd vs analytic $u_t + uu_x - \nu u_{xx}$	max error 1.2×10^{-7}
Heat residual	exact solution	max $\lvert r \rvert$ 1.2×10^{-7}
Derivative	autograd vs finite differences	agree to 2×10^{-4}
VQC bandwidth	1-qubit sweep, max harmonic vs L	equal for $L = 1, \ldots, 6$
Error metric	relative L_2 of the exact solution	4.0×10^{-8}
Parameter matching	the four base configurations	all 121
Reproducibility	same seed, Adam-only run	bit-identical

Table 4: Correctness checks (`verify_pipeline.py`, `verify_claims.py`).

B Preliminary single-seed ablations

Before the converged study we ran one-factor ablations with Adam only (2000 iterations, 1500 collocation points, seed 0). These runs are undertrained: every Burgers model has relative L_2 between 0.42 and 0.59 (except the collapsed 2π run), and the differences below are within the seed spread we later measured. They are reported for completeness; none of the comparisons is used as evidence in the main text. The classical reference in this table is a $[2, 30, 1]$ MLP sized to each QAPINN's parameter count, which differs from the $[2, 7, 11, 1]$ network of Table 2.

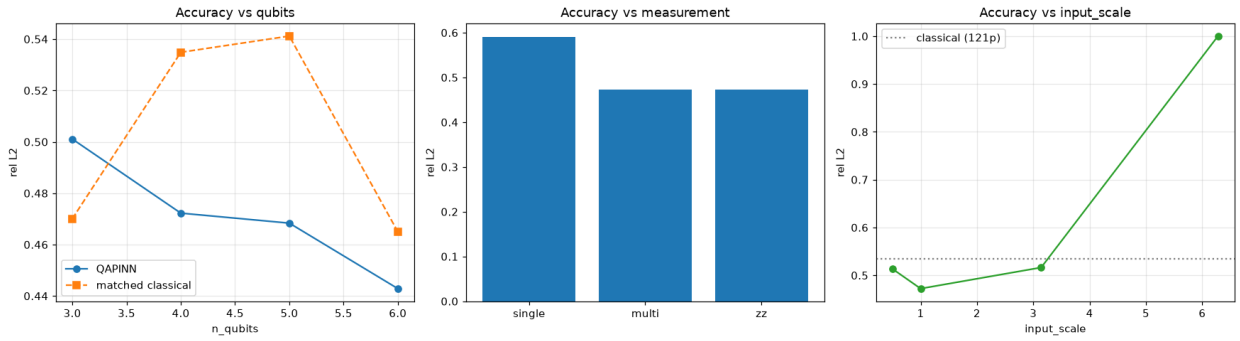


Figure 4: Preliminary sweeps over qubit count, readout and input scale (single seed, undertrained). The input-scale panel shows the collapse at 2π.

Factor	Setting	Params	QAPINN	Classical
Depth L	1	105	0.572	0.522
	2	113	0.493	0.502
	3 (base)	121	0.472	0.535
	5	137	0.424	0.500
Entanglement	none	121	0.456	0.535
	linear (base)	121	0.472	0.535
	ring	121	0.485	0.535
	all-to-all	121	0.488	0.535
Post-network	none (linear)	29	0.544	0.533
	[8]	73	0.478	0.533
	[32]	217	0.466	0.505
Qubits	3	99	0.501	0.470
	5	143	0.468	0.541
	6	165	0.443	0.465
Readout	single $\langle Z_0 \rangle$	73	0.591	0.533
	ZZ correlators	169	0.472	0.515
Input scale	0.5	121	0.513	0.535
	π	121	0.516	0.535
	2π	121	1.000	0.535

Table 5: Preliminary ablations on Burgers (relative L_2, single seed, undertrained).

Front-end comparison. Against a plain classical PINN of a different shape, the quantum front end appeared to make the post-network neurons about twice as diverse (normalised effective rank) and to increase post-network gradient variance. Against the structural control, which uses the identical post-network, neither effect was consistent: diversity was higher with the classical front end along the depth axis and with the quantum front end along the qubit axis, and the gradient-variance difference reversed at 5–6 qubits. We therefore do not claim either effect.